

Inferlane

Low-latency inference for customer-support AI agents.

Series A — August 2026

Problem

Support AI agents need sub-second responses at peak volume; general-purpose model APIs are slow and expensive at scale.

Product

A managed inference API with speculative decoding, caching and autoscaling on reserved GPUs. Drop-in replacement for the major model APIs.

Traction

Metric	Value
ARR (Aug 2026)	\$5.6M
Paying customers	44
Net revenue retention (trailing 12 months, 18 customers)	142%
Gross margin (excluding GPU inference costs, covered by cloud credits through 2027)	81%

Unit economics

Inference runs on reserved GPUs financed by \$6M of cloud credits; credits are expected to be exhausted in Q2 2027.

Team

Sofia Lindqvist, CEO — led ML platform for a large contact-center software company.

Arjun Mehta-Castell, CTO — GPU kernel engineer; 3 papers on speculative decoding.

31 people.

The round

Raising \$20M Series A, priced equity, at a \$100M pre-money valuation.

Cash \$4.0M (Aug 2026); net burn \$480k / month.